



Building Enterprise-Grade Serverless Apps

Alex Casalboni

Technical Evangelist, AWS



@alex_casalboni

About me

- Software Engineer & Web Developer
- Serverless Lover & AI Enthusiast
- AWS Customer since 2013



CUSTOMERS INNOVATING WITH AWS LAMBDA



Customer Benefits

iRobot does >1,000 Lambda deployments a day for its serverless IoT backend that runs internet connected-vacuums, with 2M connected robots by 2018 (FY17 projected).

Fannie Mae is replacing on-premises data centers with a Lambda-based solution that can run a Monte Carlo simulation on 20M mortgage calculations in 1.5 hours.

Nextdoor replaced its Apache Flume platform with a serverless data ingestion pipeline that handles 3B events daily.

HomeAway uses Lambda to process and prepare 6M user-uploaded photos a month for its vacation rental marketplace.

Agero's accident detection and driver behavior analysis platform handles over 1B Lambda requests each month and scales to handle 20x at peak load.

Revvell reduced video transcoding time by >95% at a fraction of the cost of transcoding videos on server-based solutions.

Key Trends




The Rise of Managed Services



Adoption of microservices




Stream processing and
“embarrassingly parallel” computing




Event-driven architectures

Lambda Means...



No Server Management



Flexible Scaling



High Availability




No Idle Capacity

Agenda: 6 ways to improve your serverless mojo

1. Think in patterns: the serverless architecture playbook
2. Use APIs effectively
3. SecOps like you mean it: permissions & privileges
4. Software lifecycle: from coding to deployment
5. Think in apps
6. When things go wrong: Error-handling and diagnostics

1. Serverless Patterns: Managed services + Lambda functions



Managed services as building blocks

Compute



AWS Lambda

Storage



Amazon S3

Database



Amazon DynamoDB

Amazon Aurora
Serverless

API Proxy



Amazon API Gateway

Messaging



Amazon SQS



Amazon SNS

Analytics



Amazon Kinesis



Amazon Athena

Orchestration



AWS Step Functions

Monitoring and Debugging



AWS X-Ray

Edge Compute




AWS Greengrass



Lambda@Edge

Example: Serverless analytics



Example: Serverless web app

Compute



AWS Lambda

Static Content



Amazon S3

Dynamic Content



Amazon DynamoDB



Amazon Aurora
Serverless (coming soon)

API Serving



Amazon API Gateway

Patterns for the Cloud Era

- Media transform on upload: Amazon S3 event + Lambda
- NoSQL data cleansing: Amazon DynamoDB change streams + Lambda
- Serverless website: Amazon S3 + Amazon DynamoDB + Amazon API Gateway + Lambda
- Click-stream analytics: Amazon Kinesis Data Firehose + Lambda
- Ordered event processing: Kinesis + Lambda
- Multi-function fanout: Amazon SNS (or Lambda) + Lambda
- Workflows: AWS Step Functions + Lambda
- Event distribution: Amazon CloudWatch Events + Lambda
- Serverless cron jobs: CloudWatch timer events + Lambda
- GraphQL actions: AWS AppSync + Lambda
- On-the-fly image resizing: AWS Lambda@Edge + Amazon CloudFront
- Email rules: Amazon SES + Lambda
- Configuration policy enforcement: AWS Config + Lambda
- Stored procedures: Amazon Aurora + Lambda
- Custom authorizers for APIs: API Gateway auth + Lambda
- DevOps choreography: CloudWatch alarms + Lambda
- Alexa skills: Amazon Alexa + Lambda
- Chatbots: Slack + Amazon Lex + Lambda
- IoT automation: AWS IoT + Lambda
- Smart devices: AWS Greengrass + Lambda
- On-prem file encrypt for transit: AWS Snowball Edge + Lambda
- ...

Meta-patterns

1. Service pushes async event to Lambda (S3, SNS, SQS)
2. Lambda grabs event from service (DynamoDB, Kinesis)
3. Synchronous exchange (Alexa, Lex)
4. Batch transform (Kinesis Firehose)
5. Microservice (API + Lambda + your choice of DB)
6. Customization via functions (AWS Config, SES rules)
7. Data-driven fanout (S3-Lambda, Lambda-lambda)
8. Choreography (Step Functions + Lambda)
9. Lambda functions in devices (Greengrass, Snowball Edge)

Patterns: Best practices

- Don't reinvent the wheel – use managed services when possible
 - ...and stay current 😊
- Keep events inside AWS services for as long as possible:
 - Example: S3 → Lambda → client is preferable to S3 → client → Lambda
 - Verify limits end-to-end
- Prefer idempotent, stateless functions, and when you can't...
 - Use Step Functions if you need stateful control (retries, long-running)
- Building a SaaS or ISV platform? Use the techniques AWS uses:
 1. Architect it serverlessly using Lambda and API Gateway
 2. Offer your ecosystem customization via functions
 3. Enable customers and partners to build reactive systems by emitting events

2. Use APIs Effectively



Start simple!



AWS SAM

```
AWSTemplateFormatVersion: '2010-09-09'
Transform: AWS::Serverless-2016-10-31
Resources:
  GetFunction:
    Type: AWS::Serverless::Function
    Properties:
      Handler: index.get
      Runtime: nodejs6.10
      CodeUri: s3://bucket/api_backend.zip
      Policies: AmazonDynamoDBReadOnlyAccess
      Environment:
        Variables:
          TABLE_NAME: !Ref Table
  Events:
    GetResource:
      Type: Api
      Properties:
        Path: /resource/{resourceId}
        Method: get
```


Extend with Swagger

```
AWSTemplateFormatVersion: '2010-09-09'  
Transform: AWS::Serverless-2016-10-31  
Resources:  
  Api:  
    Type: AWS::Serverless::Api  
    Properties:  
      StageName: prod  
      DefinitionUri: swagger.yaml
```

```
---  
swagger: "2.0"  
info:  
  version: "2018-06-06T01:36:07Z"  
  title: "PetStore"  
host: "0ftv5igkjd.execute-api.us-east-1.amazonaws.com"  
basePath: "/test"  
schemes:  
- "https"  
paths:  
  /:  
    get:  
      consumes:  
      - "application/json"  
      produces:  
      - "text/html"  
      responses:  
        200:  
          description: "200 response"  
          headers:  
            Content-Type:  
              type: "string"
```

Choose the right API Endpoint Type

API Endpoint Type determines where clients connect to the published APIs

- **Edge optimized:** Designed for mobile clients, uses integrated CDN (CloudFront) to enable clients to connect to the closest POP
- **Regional:** Designed for EC2 clients, within the same region as the API
- **Private:** Designed for VPC only access, with no Internet access needed. (Uses PrivateLink VPC Endpoint)

Manage throttling and rate-limiting

- Protect backends with method level rate-limiting
- Protect clients from aggressive use with Usage Plans
 - Limit on RPS
 - Limit on requests per day/week/month

New!

Create Usage Plan

Usage Plans help you meter API usage. With Usage Plans, you can enforce a throttling and quota limit on each API key. Throttling limits define the maximum number of requests per second available to each key. Quota limits define the number of requests each API key is allowed to make over a period.

Name*

Silver

Description

Usage plan for Silver-level users.

Throttling

Enable throttling

☒

Rate*

50

requests per second

Burst*

500

requests

Quota

Enable quota

☒

20000

requests per

Month


* Required

Next

ICYMI – Summary of recent API Gateway launches

- Fine-grained (method-level) throttling
- Overrides on request/response parameters and response status
- Higher limits on APIs and resource change rates
 - 600 APIs (regional or private endpoint)
 - 120 APIs (edge-optimized)
- Resource policies and cross-account access enabled
- AWS X-Ray Support (distributed tracing)
- OpenAPI 3.0 Support (import/export)

3. Security and AuthZ

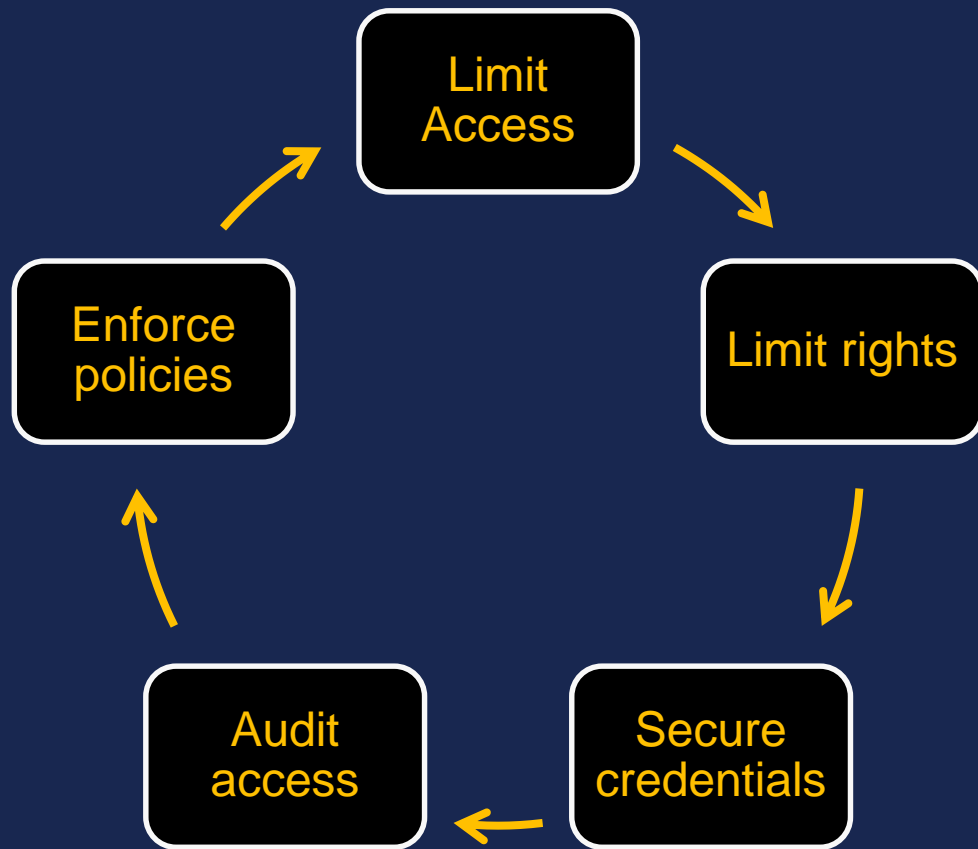


Serverless apps can be the *most* secure apps, if...

...you use the tools provided:

- Resource policies
- Custom authorizers
- Execution roles
- AWS Config policies
- AWS CloudTrail traces

Full circle of protection



Limit access: Resource policies and custom authorizers

Analogy: Who would I invite into my house?



API Gateway: Resource policies vs. Custom Authorizers

- Which should I use?
 - If resource policies can express the restriction, use them. (They're free, fast, and built in.)
- How do these two features interact?
 - We evaluate in three phases; all 3 have to pass:
 - Pre-authN (time checks, source IP checks)
 - Post-authN (call principal checks)
 - Custom authorizer (if present)
 - *We stop as soon as permission is denied, so we only run your custom authorizer if it's necessary to make a decision.*

Limits rights: Execution and invocation roles

Analogy: What do I let my kids play with?



Keeping track: Auditing and Config Policies

- Use AWS CloudTrail to track API and function management (creation, deletion, retrieval, etc.)
 - Can also audit function *invocation* for Lambda
- Run queries against CloudTrail data using Amazon Athena
 - Or build a simple analysis engine using Lambda functions
- Use AWS Config to continuously monitor and enforce policies
 - For example: Enforcing resource-level scoping on Lambda functions or custom authorizers for APIs

Calling Third Party Services Securely

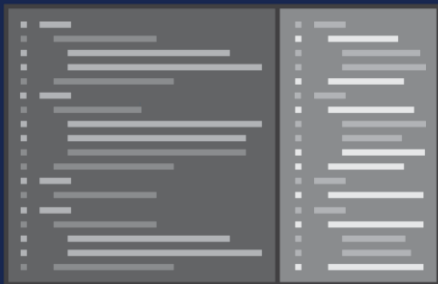

AWS Secrets Manager

- Securely store, retrieve, and manage database credentials, API keys, and other secrets for AWS or third-party systems
- Rotate, audit, and version secrets
- Built-in support for MySQL, PostgreSQL, and Amazon Aurora on Amazon RDS
- Supports both default and custom KMS keys
- *Avoid cleartext secrets in Lambda env vars or code!*

4. Building, Testing, and Deploying Serverless Apps



AWS Code Suite



AWS Cloud9
Editor

AWS
CodeStar




Best Practices for Serverless Development

- CI/CD pipeline for apps (including repo publishing!)
- Test and debug locally
- AWS CloudFormation can deploy to multiple regions
 - Useful for multi-region APIs
- Deploy incrementally for safety
 - Weighted aliases in Lambda
 - Canary stages in API Gateway
 - AWS CodeDeploy can automate SAM app deploys (including rollback)


Local test, build, & debug with SAM Local

- SAM = Serverless Application Model
- Copy of the Lambda execution environment
- Powers local test/debug
- Models serverless applications
- Works in IDEs, command line, etc.
- Use “sam init” to create a new project
- Open source



<demo time>

5. It's All about [Serverless] Apps



AWS Serverless Application Repository

- Discover, customize, and deploy serverless apps
- Publish and share apps three ways...
 - *Within* your AWS account
 - *Across* accounts
 - With everyone (*public* access)

Growing Set of App Producers



AWS Serverless Application Repository

Customize and create directly from the Lambda console, command line, or SDK

- Based on SAM – simple model of your serverless application
- Deploys using CloudFormation (use anywhere CloudFormation works!)
- Complements GitHub (source) with deployable (built) artifacts
- Easiest way to help developers get started, whether you're in an enterprise or working on an open source project
 - Will replace blueprints in the console

<demo time>

6. When Things Go Wrong: Error-Handling and Diagnostics




API Gateway CloudWatch Metrics

Default Metrics

Included for free

Granularity: per-API stage




Detailed Metrics

CloudWatch pricing

Granularity: per-method

Can be globally enabled



Lambda CloudWatch Metrics (free)

At both function and version level:

- Invocations, Errors, Throttles, IteratorAge, DeadLetterError

At both account and function level (when limit is set):

- ConcurrentExecutions

At the account level only:

- UnreservedConcurrentExecutions

API Gateway CloudWatch Logs

API Gateway Logging

- Two levels of logging, error and info
- Optionally log method request/body content
- Set globally in stage, or override per method

API Gateway Access Logging

- Customizable format for machine parsable logs

Log Pivots

- Build metrics based on log filters
- Jump to logs that generated metrics

```
5c9-11e7-8228-318bf0a162b7) Verifying Usage Plan
5c9-11e7-8228-318bf0a162b7) API Key authorized
5c9-11e7-8228-318bf0a162b7) Usage Plan check
5c9-11e7-8228-318bf0a162b7) Starting execution
5c9-11e7-8228-318bf0a162b7) HTTP Method: GET
5c9-11e7-8228-318bf0a162b7) Method request pa
5c9-11e7-8228-318bf0a162b7) Method request qu
5c9-11e7-8228-318bf0a162b7) Method request he
5c9-11e7-8228-318bf0a162b7) Method request bo
5c9-11e7-8228-318bf0a162b7) Endpoint request L
5c9-11e7-8228-318bf0a162b7) Endpoint request h
5c9-11e7-8228-318bf0a162b7) Endpoint request b
5c9-11e7-8228-318bf0a162b7) Sending request to
5c9-11e7-8228-318bf0a162b7) Received response
5c9-11e7-8228-318bf0a162b7) Endpoint response
5c9-11e7-8228-318bf0a162b7) Endpoint response
5c9-11e7-8228-318bf0a162b7) Method response b
5c9-11e7-8228-318bf0a162b7) Method response h
5c9-11e7-8228-318bf0a162b7) Successfully comp
5c9-11e7-8228-318bf0a162b7) Method completed
```

Lambda CloudWatch Logs

Basics:

- Automatically configured
- Basic info (requests, duration, memory consumption, etc.) handled automatically
- Application can add log entries easily

Cool Stuff:

- Grab a range on the console graph and jump to the logs for that time range
- Send logs to another Lambda function
- Use Elasticsearch & Kibana to aggregate/analyze

```
5c9-11e7-8228-318bf0a162b7) Verifying Usage Plan
5c9-11e7-8228-318bf0a162b7) API Key authorized
5c9-11e7-8228-318bf0a162b7) Usage Plan check
5c9-11e7-8228-318bf0a162b7) Starting execution
5c9-11e7-8228-318bf0a162b7) HTTP Method: GET
5c9-11e7-8228-318bf0a162b7) Method request pa
5c9-11e7-8228-318bf0a162b7) Method request qu
5c9-11e7-8228-318bf0a162b7) Method request he
5c9-11e7-8228-318bf0a162b7) Method request bo
5c9-11e7-8228-318bf0a162b7) Endpoint request U
5c9-11e7-8228-318bf0a162b7) Endpoint request h
5c9-11e7-8228-318bf0a162b7) Endpoint request b
5c9-11e7-8228-318bf0a162b7) Sending request to
5c9-11e7-8228-318bf0a162b7) Received response
5c9-11e7-8228-318bf0a162b7) Endpoint response
5c9-11e7-8228-318bf0a162b7) Endpoint response
5c9-11e7-8228-318bf0a162b7) Method response b
5c9-11e7-8228-318bf0a162b7) Method response h
5c9-11e7-8228-318bf0a162b7) Successfully comp
5c9-11e7-8228-318bf0a162b7) Method completed
```

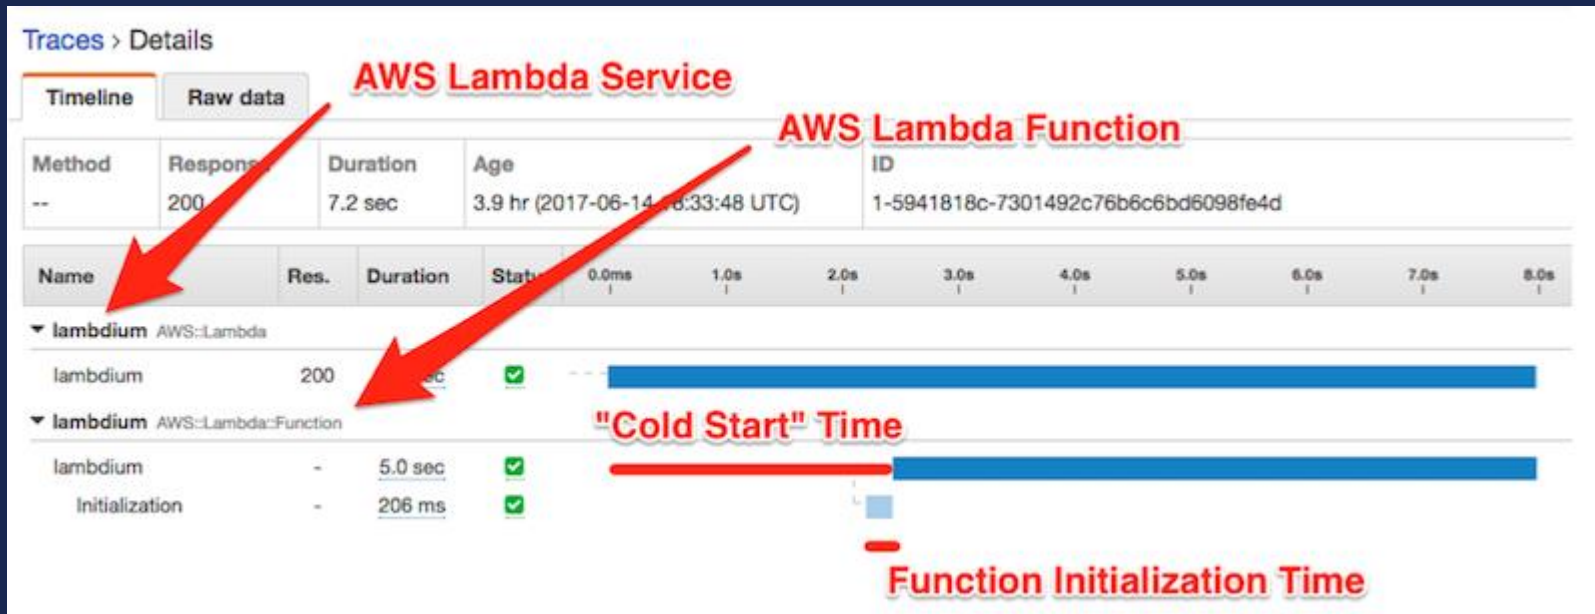
The Lambda Function Request Lifecycle

1. Download your code (as a SquashFS filesystem)
2. Start the language runtime
3. Load your code (e.g., class loading in Java)
4. Run “init” code (e.g., static initializers in Java)
5. Process the request

} **warm**

} **cold**

Same View in AWS X-Ray



Things You Can Track in X-Ray for Lambda

- Cold vs warm starts
- Async dwell time (time event spends in queue)
- Duration of calls to other AWS services (and with a little work to add trace points, third-party services)
- Errors and throttles from AWS service calls

Lots of great third-party offerings can also help!

Making Cold Starts More Efficient

- Control the dependencies in your function's deployment package
 - Tree shake your Java code! (Hint: Spring is expensive 😞)
- Can combine 'cold' and 'hot' functions
- Optimize for your language, where applicable...
 - Node – Browserfy, Minify
 - AWS Java v2 SDK (preview)
- Only use VPC if you need access to a resource there!
- Delay loading where possible

Summary: 6 ways to improve your serverless mojo

1. Think in patterns: the serverless architecture playbook
2. Use APIs effectively
3. SecOps like you mean it: permissions & privileges
4. Software lifecycle: from coding to deployment
5. Think in apps
6. When things go wrong: Error-handling and diagnostics



Thank you!

Alex Casalboni

Technical Evangelist, AWS



@alex_casalboni